

A Lesson on Instruction Set Architectures (ISAs)

Application software

*Systems software
(OS, compiler)*

— ISA —

*Digital design and
architecture*

Circuits and devices


Andrew DeOrio

awdeorio@umich.edu

andrewdeorio.com



ARM



Today's lesson

- Based on material from my sophomore/junior level Introduction to Computer Organization course (EECS 370 at Michigan)
- Slides online at andrewdeorio.com
- In previous lectures, we covered:
 - Basic programming in C
 - Representing numbers in binary and hexadecimal
 - Computer history

Review: programming in C

- Program

```
/* hello.c */  
#include <stdio.h>  
  
int main() {  
    printf("hello world!\n");  
    return 0;  
}
```

- Compile

```
$ gcc hello.c
```

- Run


```
$ ./a.out  
hello world!
```

Review: numbers in binary


- 42 in base 10 (AKA decimal)
 $= 40 + 2$
 $= 4 * 10^1 + 2 * 10^0$
 $= 42_{10}$
- 42 in base 2 (AKA binary)
 $= 32 + 0 + 8 + 0 + 2 + 0$
 $= 1 * 2^5 + 0 * 2^4 + 1 * 2^3 + 0 * 2^2 + 1 * 2^1 + 0 * 2^0$
 $= 101010_2$
- 42 in base 16 (AKA hexadecimal)
 $= 32 + 10$
 $= 2 * 16^1 + 10 * 16^0$
 $= 2A_{16}$
or
 $= 0010\ 1010_2$
 $= \quad 2 \quad \quad A$
 $= 0x2A$

Decimal	Binary	Hex
0	0000	0x0
1	0001	0x1
2	0010	0x2
3	0011	0x3
4	0100	0x4
5	0101	0x5
6	0110	0x6
7	0111	0x7
8	1000	0x8
9	1001	0x9
10	1010	0xA
11	1011	0xB
12	1100	0xC
13	1101	0xD
14	1110	0xE
15	1111	0xF

Review: computer history




Astrolavos computed phases of the moon




Babbage's analytical engine. Mechanical general purpose computer.



ENIAC electronic computer, using vacuum tubes




1970: Intel 4004, first microprocessor, 2,250 transistors



2014: Intel Core i7 (Broadwell), 1.4B transistors

Moore's law



iPhone 6 teardown




Photo credit: ifixit.com



A8: 2B
transistors

Abstraction

- How do you put together 2B transistors?
- Use abstraction to simplify design
- Separate lower-level details from higher-level details
- The abstraction we'll focus on today is called the *Instruction Set Architecture*, or ISA
- Put another way, we need a language that lets SW talk to HW

Abstraction



Application software

Systems software (OS, compiler)

————— Instruction Set Architecture (ISA) —————

Digital design and architecture

Circuits and devices



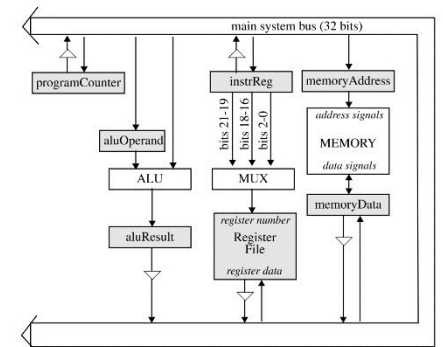
Abstraction

- Alice is a hardware engineer, Bob is a software engineer
- Alice and Bob agree on an ISA
- Alice creates a chip
 - Implements the ISA
- Bob writes software
 - Uses the ISA
- They put it together and the iPhone 6 works



ISA examples

- Core i7 implements the x86 ISA
- Apple A8 implements ARM ISA
- The Little Computer implements the LC2K ISA
 - Example ISA for this class



RISC and CISC

- How to encode instructions?
- CISC: Complex Instruction Set Computer
 - Instructions can vary in size
 - Example: x86 (most laptops)
- RISC: Reduced Instruction Set Computer
 - All instructions are same length
 - Example: ARM (many cell phones)



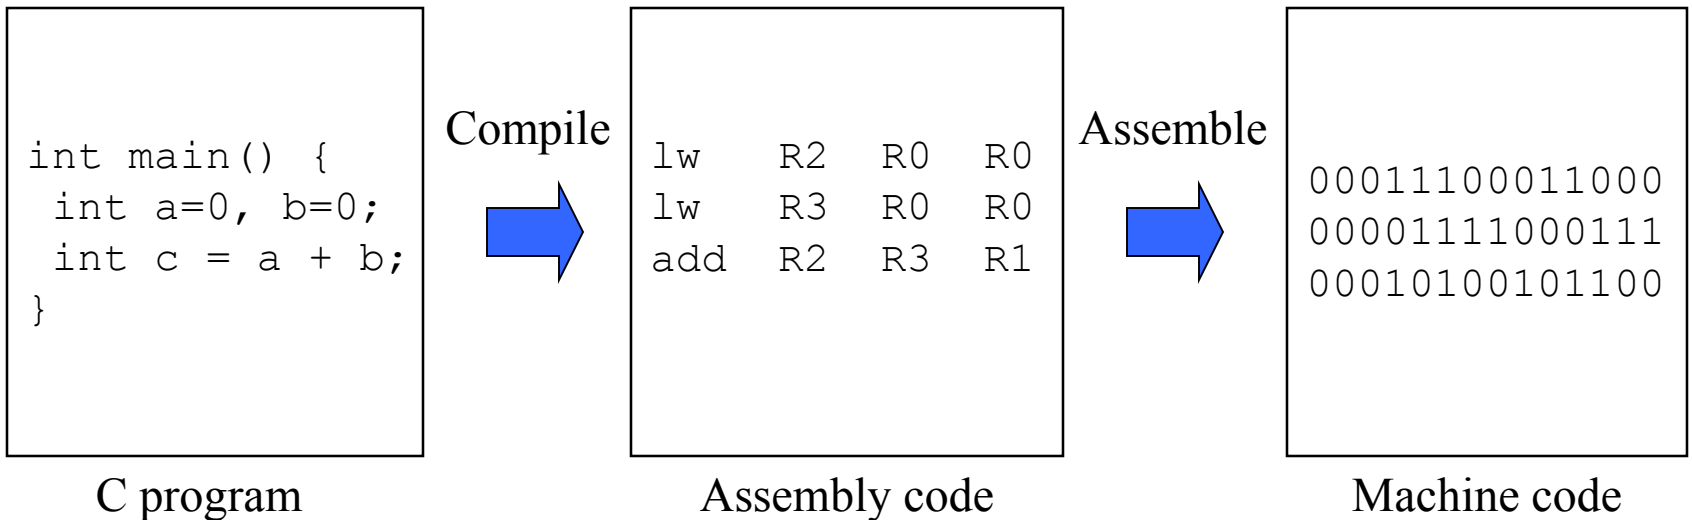
Instructions

- The language that hardware and software use to talk to each other is the ISA
- The vocabulary consists of *instructions*
- Examples of *instructions*
 - Add two numbers
add R2 R3 R1
 - Check to see if two numbers are equal
beq R4 R5 5
 - Copy a piece of data
lw R6 R7 5

Software program to instructions

A program that converts
input C-programs to
output assembly programs
Example: `gcc -S`

A program that converts
input assembly programs to
output binary programs
Example: `as` or (part of) `gcc`



Storing instructions


- First, break up the instruction into fields
- Opcode – what instruction to perform
- Source (input) operands
- Destination (output) operand

opcode	src1	src2	dest
add	R2	R3	R1


Translation: value in R1 = contents of R2 + R3

Storing operands

- Operands are stored in registers
`add R2 R3 R1`
- A register is a physical piece of hardware that stores one number
- `add` is performed by physical hardware on chip
- Registers are small and fast
- In future lessons, we'll learn how main memory extends this storage




Assembly instruction encoding



- m bits can encode 2^m different values
- n values can be encoded in $\lceil \log_2(n) \rceil$ bits


Exercise

- What is the maximum number of registers that can be designed in an ISA with:
 - 16-bit instructions
 - 100 opcodes
 - All instructions have 2 source operands and 1 destination operand



Solution

- What is the maximum number of registers that can be designed in an ISA with:
 - 16-bit instructions
 - 100 opcodes
 - All instructions have 2 source operands and 1 destination operand



1. num opcode bits = $\lceil \log_2(100) \rceil = 7$
2. num bits for operands = $16 - 7 = 9$
3. num bits per operand = $9 / 3 = 3$
4. maximum number of registers = $2^3 = 8$

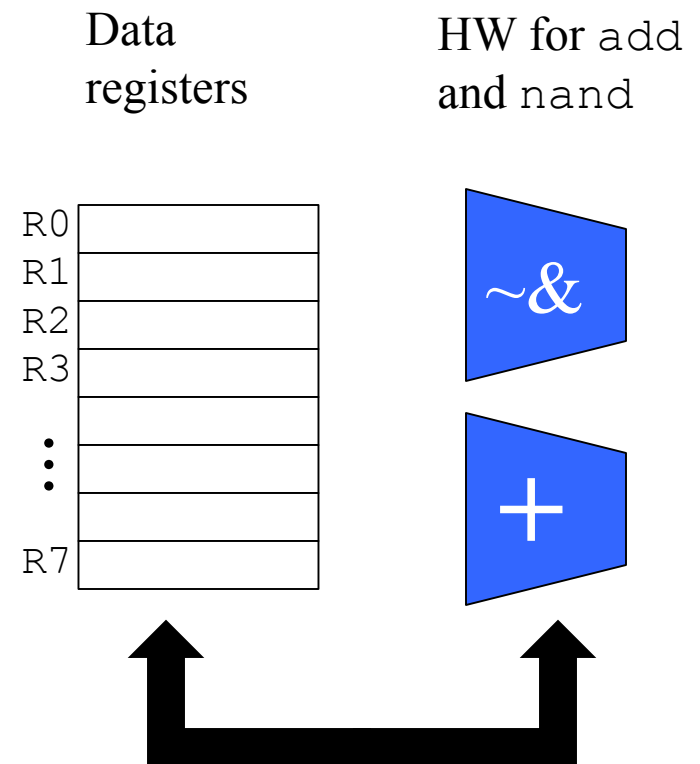
News break

- Let's take a break
- What's going on in the computer-related news?

LC2K ISA

- Little Computer 2000 (LC2K) is a small example computer and Instruction Set Architecture
- 32-bit processor
 - Instructions are 32 bits
 - Registers are 32 bits
- 8 registers
- Supports 65,536 ($=2^{16}$) storage locations (memory addresses)
- 8 instructions
 - `add`, `nand`, `lw`, `sw`, `beq`, `jalr`, `halt`, `noop`

A portion of the
LC2K hardware



LC2K assembly instructions

Instruction	Example	Meaning in code	Meaning in English
add	add 1 2 3	$R3 = R1 + R2$	Add two numbers
nand	nand 1 2 3	$R3 = \sim (R1 \ \& \ R2)$	Negative <i>and</i>
lw	Instructions for copying data. We'll learn about these next time.		
sw			
beq	Control flow instructions used by loops, if-statements and functions. Also a future lesson.		
jalr			
halt	halt		Stop the program
noop	noop		Do nothing

Storing operands

- Example program
- Assume all registers are initialized to zero
- Remember: last register is the destination

```
nand 1 1 2
```

```
add 1 2 3
```

R0	0
R1	0
R2	0
R3	0
	0
⋮	0
	0
R7	0

Storing operands solution

- Example program
- Assume all registers are initialized to zero
- Remember: last register is the destination

```
nand 1 1 2
add 1 2 3
```

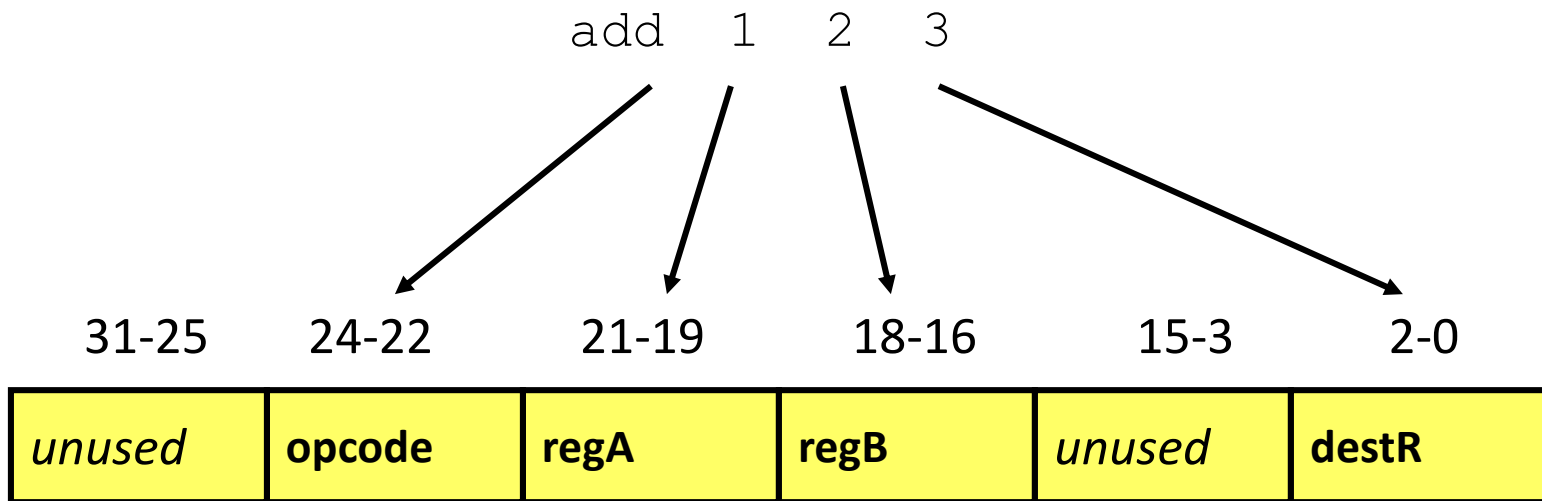
R0	0
R1	0
R2	0xFFFF FFFF
R3	0xFFFF FFFF
	0
⋮	0
	0
R7	0

Assembly instruction encoding

- Instructions are stored in the same way as data
- Instructions are encoded as numbers
- How to encode it? Use the Architecture Reference Manual
- Excerpt from the LC2K manual
- `add srcA srcB dest`
 - bits 24-22: opcode 000
 - bits 21-19: src A operand
 - bits 18-16: src B operand
 - bits 15-3: unused (0's)
 - bits 2-0: dest operand


Instruction encoding

- Instructions are stored in the same way as data
- Instructions are encoded as numbers
- How to encode it? Use the Architecture Reference Manual



Instruction formats


- Positional organization of bits in LC2K
 - Implies nothing about bit values!!!
- Both add and nand are encoded this way



- Opcode encodings
 - add (000), nand (001)
- Register values
 - Just encode the register number (R2 = 010)

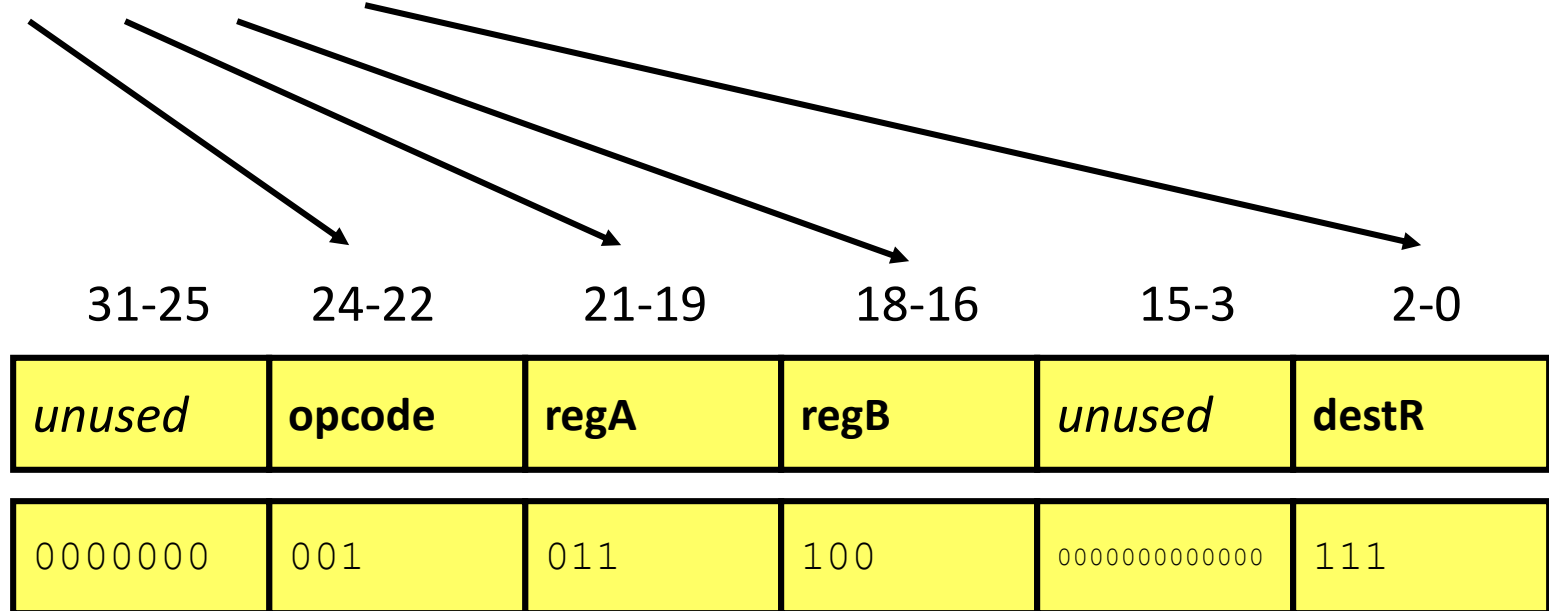
Encoding exercise

- `nand 3 4 7 //R7 = ~(R3 & R4)`
- `nand` opcode is 001



Encoding exercise solution

- `nand 3 4 7 //R7 = ~(R3 & R4)`



bin: 0000000 001 011 100 0000000000000000 111

bin: 0000 0000 0101 1100 0000 0000 0000 0111

hex 0 0 5 C 0 0 0 7

hex: 0x005C0007

Recall our abstraction

- Hardware engineer Alice and software engineer Bob
- Alice and Bob agree on an ISA
- Alice creates a chip
 - Implements the ISA
- Bob writes software
 - Uses the ISA
- They put it together and the iPhone 6 works



Abstraction exercise

- For each of these scenarios, does the software need to change?
Does the hardware need to change?
- Change `nand` instruction to `and`
- Add a new `mul` (multiply) instruction
- New piece of hardware that can add faster
- New operating system that is better at multitasking

Abstraction exercise solution

- For each of these scenarios, does the software need to change?
Does the hardware need to change?
- Change `nand` instruction to `and`
 - Both software and hardware
- Add a new `mul` (multiply) instruction
 - Hardware changes. Software optional.
- New piece of hardware that can add faster
 - Hardware only
- New operating system that is better at multitasking
 - Software only



The power of ISAs

- We can upgrade the hardware and the old software still works!
 - Example: iPhone 5 to iPhone 5s you can still run your old applications.
- We can upgrade the software and the old hardware still works!
 - Example: iOS7 to iOS8 upgrade on iPhone 5/5s. You can still use your old hardware.



Next time

- Storage and memory
 - How do the load and store instructions work?
- Control instructions
 - How does the `beq` instruction work?
 - How do we implement loops and `if`-statements?

Programming assignment #1

- Coming soon
- Write an assembler to convert input (assembly language program) to output (machine code version of program)
- Write a behavioral simulator to run the machine code version of the program (printing the contents of the registers and memory after each instruction executes)
- Write an efficient assembly language program to multiply two numbers